



Dr. APJ Abdul Kalam School of Engineering

# ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING

## You Only Look Once (YOLO)

Presented By: Kiran T  
23BECS132

FACULTY: Prof. Gnanasundari

# Introduction to Object Detection

Object detection is a computer vision task in Artificial Intelligence that identifies and locates objects within an image or video. Unlike simple image classification, which only tells *what* object is present, object detection also specifies *where* the object is by drawing bounding boxes around it.

It is widely used in applications such as:

- **Self-driving cars** – detecting pedestrians, vehicles, and traffic signs.
- **Surveillance systems** – recognizing suspicious activities in real time.
- **Healthcare** – identifying tumors or anomalies in scans.

# Introduction to YOLO

- YOLO (**You Only Look Once**) is a deep learning algorithm for real-time object detection, introduced by Joseph Redmon in 2016.
- Unlike older methods that scan an image many times, YOLO processes the entire image in one pass, making it extremely fast and efficient.

## **Key Highlights:**

- Detects multiple objects at once
- Treats detection as a single regression problem
- Combines speed and accuracy, ideal for real-world use



# Importance in AIML

- YOLO is an excellent example of how Artificial Intelligence, Machine Learning, and Deep Learning come together to solve real-world problems.
- **AI:** Enables machines to “see” and interpret visual data.
- **ML:** Learns from large datasets to classify and locate objects.
- **DL (CNNs):** Extracts patterns and features directly from images.

Because of its speed and accuracy, YOLO is widely used in AIML projects such as autonomous driving, healthcare imaging, and security systems.

# Why YOLO?

- Traditional object detection methods like R-CNN and Fast R-CNN were accurate but too slow for real-time use. They relied on multiple stages such as region proposals, feature extraction, and classification.

YOLO solved this by:

- Performing detection in a single step.
- Looking at the entire image at once instead of scanning parts.
- Offering a balance of speed and accuracy.

# YOLO Overview

- YOLO changes object detection by treating it as a single regression task from images to bounding boxes and class labels.
- The image is divided into a grid.
- Each grid cell predicts bounding boxes and class probabilities.
- Predictions are combined to detect multiple objects at once.
- YOLO analyzes the entire image in one pass, it gains a global understanding of context, reducing false detections and making results faster.

# Architecture

- **Input Layer**
  - Image resized (commonly  $416 \times 416$  or  $640 \times 640$ ).
  - Normalized for faster processing.
- **Backbone (Feature Extractor)**
  - Deep CNN extracts visual features such as shapes, edges, and textures.
  - Different versions use backbones like Darknet-19, Darknet-53, or CSPDarknet.
- **Detection Head**
  - Predicts bounding box coordinates, objectness score, and class probabilities.
  - Produces final detections for multiple objects.



# How YOLO Works

YOLO detects objects in a single pass using the following steps:

- **Input Image**
  - The image is resized and divided into a grid of cells.
- **Bounding Box Prediction**
  - Each grid cell predicts multiple bounding boxes and a confidence score for each box, indicating the probability of an object being present.
- **Class Prediction**
  - Each bounding box also predicts class probabilities for possible object types (e.g., car, person, dog).
- **Non-Maximum Suppression (NMS)**
  - Overlapping boxes are filtered to keep only the most accurate detection.



# YOLO Versions

YOLO has evolved over time to improve accuracy, speed, and versatility:

- **YOLOv1 (2016):** First version; fast but struggled with localization and small objects.
- **YOLOv2 (YOLO9000, 2017):** Improved accuracy and could detect over 9000 object categories.
- **YOLOv3 (2018):** Introduced multi-scale detection for better accuracy on small and large objects.
- **YOLOv4 (2020):** Optimized for GPU; faster and more accurate.
- **YOLOv5 (2020):** Lightweight, popular in industry; easy to train.
- **YOLOv8 (latest):** Lightweight, edge-device friendly, high accuracy, supports PyTorch & TensorFlow.

Each version improves detection speed, accuracy, and adaptability, making YOLO suitable for real-time AIML applications.

# Key Features

YOLO stands out in object detection due to its combination of speed, accuracy, and simplicity:

- **Real-Time Detection:** Can process images at over 45 frames per second.
- **Single-Pass Detection:** Predicts bounding boxes and class probabilities in one go.
- **End-to-End Trainable:** No separate pipelines for feature extraction or region proposals.
- **Versatile:** Works across a wide range of environments and object types.
- **Global Understanding:** Looks at the entire image, reducing false positives compared to region-based methods.

# YOLO vs Other Models

| Feature           | YOLO                   | R-CNN / Fast R-CNN                 |
|-------------------|------------------------|------------------------------------|
| <b>Speed</b>      | Very fast (real-time)  | Slow, multi-stage process          |
| <b>Detection</b>  | Single-pass prediction | Requires region proposals          |
| <b>Training</b>   | End-to-end trainable   | Multi-stage training needed        |
| <b>Accuracy</b>   | High (newer versions)  | Slightly higher for complex images |
| <b>Complexity</b> | Simple architecture    | More complex, slower               |

YOLO achieves a balance of speed and accuracy, making it suitable for real-time applications, while R-CNNs may be more accurate but too slow for live tasks.



# Applications

YOLO's speed and accuracy make it ideal for a wide range of real-world applications in AIML:

- **Autonomous Vehicles:** Detects pedestrians, vehicles, traffic signs, and road obstacles in real time.
- **Surveillance Systems:** Monitors public spaces for suspicious activity or intruders efficiently.
- **Healthcare:** Helps identify tumors, lesions, or other anomalies in medical imaging.
- **Retail & Industry:** Automates inventory management, counts items, and tracks movement.
- **Robotics:** Guides robots to recognize and manipulate objects accurately.
- **Agriculture:** Monitors crops, detects pests, and tracks livestock for smart farming solutions.

YOLO is widely used across industries because it transforms visual data into actionable intelligence quickly and reliably.



# Advantages

- **Real-Time Performance:** Processes images at high speeds, suitable for live video and real-time applications.
- **Multiple Object Detection:** Can detect and classify several objects in a single image simultaneously.
- **Unified Architecture:** Combines feature extraction, classification, and bounding box prediction in one network.
- **Reduced False Positives:** Global image analysis helps minimize incorrect detections compared to region-based methods.
- **Versatility:** Works effectively across various domains such as automotive, healthcare, security, and robotics.

# Limitations

While YOLO is fast and efficient, it has some limitations that should be considered:

- **Difficulty with Small Objects:** YOLO may miss small or closely packed objects due to its grid-based detection approach.
- **Localization Errors:** Fixed grid divisions can sometimes cause bounding boxes to be slightly inaccurate.
- **Lower Accuracy in Complex Scenes:** For images with many overlapping objects, YOLO may be less precise than some region-based detectors.
- **Data Requirements:** Requires large, annotated datasets for training to achieve high accuracy.

Despite these limitations, newer versions like YOLOv5 and YOLOv8 have addressed many issues, improving both accuracy and flexibility for real-world applications.

# Conclusion

- YOLO (You Only Look Once) is a real-time object detection algorithm.
- Combines speed, accuracy, and simplicity in a single neural network.
- Can detect multiple objects simultaneously in images and videos.
- Widely used in AIML applications: autonomous vehicles, healthcare, surveillance, retail, and agriculture.
- Continuous improvements (YOLOv5, YOLOv8) make it more accurate and lightweight.
- Demonstrates how AI, ML, and DL work together to turn visual data into actionable intelligence.